

# CONCURRENCY

## SOME BASICS

Joseph Kehoe<sup>1</sup>

<sup>1</sup>Department of Computing and Networking  
Institute of Technology Carlow

CDD101, 2017

# TABLE OF CONTENTS

- 1 INTRODUCTION
- 2 STRATEGIES
- 3 CATEGORIES OF PARALLELISM
- 4 ARCHITECTURES
- 5 PERFORMANCE

# TABLE OF CONTENTS

- 1 INTRODUCTION
- 2 STRATEGIES
- 3 CATEGORIES OF PARALLELISM
- 4 ARCHITECTURES
- 5 PERFORMANCE

# COMPOSITION OF CONCURRENT ALGORITHMS

- Algorithms consist of one or more tasks acting on data
- Dependencies will exist between the tasks

**DATA DEPENDENCY** one task requires data to be “prepared” by another task before it can start

**CONTROL DEPENDENCY** Task side effects need to be ordered e.g. I/O Operations

**Fork-Join** is one popular way of managing these dependencies

- New control flows (concurrent tasks) are created at a fork point
  - One splits into many
- Synchronisation occurs when tasks are merged into a single control flow
  - Many become one
- Each control flow is sequential

# TABLE OF CONTENTS

- 1 INTRODUCTION
- 2 STRATEGIES
- 3 CATEGORIES OF PARALLELISM
- 4 ARCHITECTURES
- 5 PERFORMANCE

## Data Parallelism

- Parallelism grows as data grows
- This is scalable

## Functional Parallelism

- Divide task into multiple concurrent tasks
- Not as scalable (why?)

# REGULAR VERSUS IRREGULAR

## Regular Parallelism

- Tasks are similar
- Tasks have predictable dependencies
- E.g. Matrix Multiplication

## Irregular Parallelism

- Tasks are dissimilar in a manner that creates unpredictable dependencies
- E.g. Search in games



# TABLE OF CONTENTS

- 1 INTRODUCTION
- 2 STRATEGIES
- 3 CATEGORIES OF PARALLELISM
- 4 ARCHITECTURES
- 5 PERFORMANCE

# THREAD PARALLELISM

- Each task has its own flow of control
- Useful for all types of parallelism

**HARDWARE THREAD** Thread that is supported by hardware e.g. separate core for each thread (parallel)

**SOFTWARE THREAD** Software based thread e.g. time slicing of processor time (concurrent)

**HYPERTHREAD** Core has duplicated some components to allow it to run two threads at once

# VECTOR PARALLELISM

- Single control flow can operate on multiple data elements
  - Useful for regular parallelism (mainly)
- Intel AVX allows register to hold many (8) 32 bit floating point numbers
  - All can be acted on simultaneously
- Requires less silicon to implement than thread based parallelism
- This approach can emulate thread parallelism by using Packing or Masking
  - These “threads” are called fibers

# TABLE OF CONTENTS

- 1 INTRODUCTION
- 2 STRATEGIES
- 3 CATEGORIES OF PARALLELISM
- 4 ARCHITECTURES
- 5 PERFORMANCE

# FLYNN'S CATEGORIES

**SISD** Single Instruction, Single Data

**SIMD** Single Instruction, Multiple Data

**MIMD** Multiple Instruction, Multiple Data

**MISD** Multiple Instruction, Single Data

GPU Vendors use:

**SIMT** Single Instruction, Multiple Threads (SIMT). A Tiled SIMD where each SIMD processor emulates multiple threads (fibers, really) using masking

There is a memory hierarchy where each level can be more than an order of magnitude slower than the next (from fastest to slowest)

- Registers (on each core)
- L1 Cache (Instruction and data caches) on each core
- L2 Cache shared between multiple cores
- L3 Cache one per processor
- RAM Shared by everyone on board
- Main Memory (SSD or mechanical) Shared by everyone on box

# TABLE OF CONTENTS

- 1 INTRODUCTION
- 2 STRATEGIES
- 3 CATEGORIES OF PARALLELISM
- 4 ARCHITECTURES
- 5 PERFORMANCE

Lots of things affect performance but at a high level remember the following:

**DATA LOCALITY** Keep data close to the thread using it

**SLACK** Have more potential tasks than there is actual parallelism

**AVOID HAVING TOO MANY THREADS** One per core is ideal. Use a thread pool!

More on performance later!